

How to Build the Course Recommender — Step by Step

This is your **construction manual**. The Handbook told you *what* the project is and *why*. This document tells you *exactly how to build it* — module by module, what goes in, what comes out, and how each piece connects to the next inside one Colab notebook.

[9 modules · 1 notebook](#)[read top → bottom](#)[companion to the Handbook](#)

00

How to use this document

You are building **one Jupyter notebook in Google Colab**. Think of it as a factory line: data enters at the top, and each station (a **module**) does one job and hands its result to the next station. A module is just **one labelled section of the notebook** — usually 1–4 code cells.

The golden rule of this project: the *output* of one module is the *input* of the next. The connection between them is always a **named Python variable**. Build the modules **in order** and keep the variable names exactly as written here, and each step will feed cleanly into the next.

How to work through it

1. Read the [Pipeline diagram](#) once, fully. That is the whole project on one screen.
2. Keep the [Data Dictionary](#) open in a second tab — it is the master list of every variable that travels between modules. It is your contract.
3. Build **in order**, **0 → 7**. Do not start Module 4 before Module 2 produces clean data. Module 8 is optional and last.
4. For your assigned module: read its *Purpose*, copy the *Starter code* into a cell, run it, then check the "*Done when...*" box at the bottom before moving on.
5. When something breaks, check *Common errors* first — your bug is almost certainly listed there.


The starter code is a scaffold, not the finished answer. It runs, but it is intentionally minimal. Your job — and the part that earns marks — is to understand it, improve it, and explain it in the report. Copying without understanding will show immediately when you present.

One example runs through this entire document. Meet **Riya**: she already knows `Python`, and she wants to become a `Data Scientist`. Watch her data flow through every module so you can see exactly what each step produces.

01

The full pipeline — how everything connects

Every box is a module (a section of the notebook). Every **labelled arrow is the variable** that one module hands to the next. Modules 4 and 5 are **two different methods**: they both read the same two inputs and each produces its own ranked list, which meet again at the evaluation step (M6), where you compare them.



tfidf_recs + embed_recs

6

Evaluate & compare ★ the graded heart

in: `tfidf_recs · embed_recs · is_relevant()` → out: `eval_table · charts · winner`

winner · best_recs

7

Pick best + LLM roadmap (Groq)

in: `best_recs · current_skills · target_role` → out: `roadmap_text`

roadmap_text

8

Optional UI — only if time allows

wraps Modules 3 → 7 in input boxes → out: **interactive demo**

Read it as one sentence: load the data → work out which skills are missing → let two different methods recommend courses → measure which method is better → hand the winner's picks to an LLM to write a friendly roadmap.

02

The data dictionary — your contract

These are **every variable that travels between modules**. Spell them exactly like this in your code — that consistency is what lets each module's output feed cleanly into the next as you build the notebook in order.

VARIABLE	TYPE	CREATED IN	WHAT IT HOLDS (EXAMPLE)
<code>courses_raw</code>	DataFrame	M1	The CSV exactly as downloaded, unmodified
<code>role_skills</code>	dict	M1	<code>{"Data Scientist": ["python", "statistics", "machine learning", ...]}</code>
<code>courses</code>	DataFrame	M2	Clean table. Required columns: <code>course_id</code> , <code>title</code> , <code>skills</code> , <code>text</code> (extra columns like <code>description</code> , <code>provider</code> may ride along but aren't used)
<code>current_skills</code>	list[str]	M3	<code>["python"]</code>
<code>target_role</code>	str	M3	<code>"Data Scientist"</code> — must be a key in <code>role_skills</code>
<code>skill_gap</code>	list[str]	M3	<code>["statistics", "machine learning", "sql", "data visualization", "deep learning"]</code>
<code>query_text</code>	str	M3	<code>"statistics machine learning sql data visualization deep learning"</code>
<code>tfidf_recs</code>	DataFrame	M4	Ranked list. Columns: <code>course_id</code> , <code>title</code> , <code>score</code> , <code>rank</code>
<code>embed_recs</code>	DataFrame	M5	Same shape as <code>tfidf_recs</code>
<code>is_relevant()</code>	function	M6	Decides if a course counts as a correct recommendation (see §03)
<code>eval_table</code>	DataFrame	M6	precision@k for each method, side by side
<code>winner</code>	str	M6	<code>"Embedding"</code> or <code>"TF-IDF"</code>
<code>best_recs</code>	DataFrame	M7	The winning method's ranked list (= <code>tfidf_recs</code> or <code>embed_recs</code>)
<code>roadmap_text</code>	str	M7	The month-by-month roadmap the LLM wrote

Build in order. The evaluation (M6) cannot run until both recommenders (M4, M5) produce their `_recs` tables in *exactly* this shape. Agree the column names early and never change them silently.

03

Two rules that are locked — do not redesign these

These two decisions are made *for you*, because getting them wrong quietly breaks everything downstream. Everything else is yours to design.

Rule 1 — What makes a course "relevant" (the ground truth)

To *measure* a recommender, you need a definition of a correct answer. The raw data does not contain one, so we define it:

A recommended course is *relevant* if one of the gap skills appears (as a whole word) in its skills.

i.e. the course teaches at least one skill the person is actually missing.

```
def is_relevant(course_skills, skill_gap):  
    # True if a gap skill appears as a whole word in the course's skills  
    skills_text = " ".join(course_skills).lower()  
    return any(re.search(r"\b" + re.escape(g) + r"\b", skills_text) for g in skill_gap)
```

Why not the simpler `set(course_skills) & set(skill_gap)` ? Because the real dataset's skills are *phrases* — e.g. "Statistics for Data Science", not the bare token "statistics". An exact set-match would score `statistics` as 0 relevant courses even though dozens teach it (verified on the real data). Whole-word matching (`\b...\b`) catches `statistics` inside "statistics for data science", yet won't wrongly match `sql` inside "mysql".

This is the single most important rule in the whole project — it is the yardstick every score is measured against. **Discuss its limitations in your report** (e.g. it treats all gap-skills as equally important, and whole-word matching can still miss synonyms like "ML" for "machine learning"). That discussion is worth marks.

Rule 2 — `target_role` must match a known role

The role the user picks must be a key that exists in `role_skills`. If you allow free typing, `"data scientist"` (lowercase) or `"ML Engineer"` will silently find nothing and the pipeline produces empty results with no error.

So: the user *chooses from a fixed list of roles* (the keys of `role_skills`) — never types a role freely. In the optional UI (M8) this becomes a dropdown, not a text box.

04

Weekly plan — what to build, test & demo each week

This is your **sprint plan**. The weeks between Thursday meetings are your sprints. Each sprint you **build** a set of modules, **test-run** them on real data, and **demo them working** at the next meeting. Build strictly in order — a later module can't run until the earlier one produces its output.

A module is "done" only when you have (1) run it on real data, (2) ticked its green *"Done when..."* box, and (3) can *show it working* at the next meeting. Code that "should work" but was never run does not count.

SPRINT	BY THIS MEETING	BUILD	TEST-RUN / DEMO AT THE MEETING
1	Mtg 2 · 25 Jun	M0 Setup · M1 Load data · a <i>crude</i> M4 (TF-IDF returns <i>something</i>)	Notebook opens in Colab · data loads · <code>role_skills</code> works · a rough recommendation prints
2	Mtg 3 · 2 Jul	M2 Clean · M3 Skill gap · finish M4 (clean TF-IDF) · M5 Embeddings	Both recommenders return a top-K list for the same test user
3	Mtg 4 · 9 Jul	M6 Evaluate & compare (precision@k + charts)	Comparison table + bar chart shown · you can say which method won and why
4	Mtg 5 · 16 Jul	M7 LLM roadmap · feature freeze · decide on optional M8	Full end-to-end run: input → recommendations → roadmap
5	Mtg 6 · 23 Jul	(optional M8) · clean the notebook · draft your report sections	Notebook runs top-to-bottom with no errors · each owner's draft section
6	Mtg 7 · 30 Jul	Finalise report · dry-run the walkthrough	Full assembled report draft · practice presentation
—	31 Jul (Fri)	—	Submit the final report to ISI

Timeline reality: kickoff (Fri 19 Jun) to deadline (31 Jul) is ≈ 6 working weeks. It's tight — that's why scope is lean and why you build a *crude* version first (Sprint 1) and improve it, rather than waiting for perfect. **No new features after the Sprint-4 freeze.**

05

The modules, in build order

Build these in sequence. Each one: *what it does* → *exact inputs* → *exact outputs* → *how to build it* → *Riya's worked example* → *common errors* → *done-when*.

MODULE 0

Setup & environment

1 Sprint 1 · Mtg 2

Purpose. Install the libraries and load your free Groq key, once, at the top of the notebook. Everything below depends on this cell having run.

INPUT → OUTPUT

IN	OUT
nothing	all libraries importable · <code>GROQ_API_KEY</code> available via <code>os.environ</code>

STARTER CODE

```
# run once, at the very top
!pip install -q sentence-transformers groq

import pandas as pd, numpy as np, re, os
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity
import matplotlib.pyplot as plt

# paste your free Groq key here (get it from console.groq.com – no credit card)
os.environ["GROQ_API_KEY"] = "gsk_...your key..."
```

Never commit your API key to GitHub. Before you push, replace the key with `"gsk_...your key..."` again, or use Colab's 🔑 Secrets panel. A leaked key gets disabled automatically.

📄 **For your report:** note your build environment — Google Colab, Python version, and the libraries you used (with versions: `pd.__version__` etc.). Feeds the report's *Setup / Tools* section.

✅ **Done when:** the cell runs with no red errors and `print(os.environ["GROQ_API_KEY"][:4])` prints `gsk_`.

MODULE 1 Load data

📅 Sprint 1 · Mtg 2

Purpose. Read the course catalogue into a DataFrame, and define the role→skills table. Do *nothing* else here — no cleaning yet. Just get the raw material into memory.

INPUT → OUTPUT

IN	OUT
a course CSV file (uploaded to Colab) · the role list you decide on	<code>courses_raw</code> (DataFrame) · <code>role_skills</code> (dict)

THE COURSE CSV — WHAT IT MUST CONTAIN

You're using one fixed dataset: **Coursera Courses & Skills 2024** (Kaggle), the file `coursera_course_dataset_v3.csv` — **623 rows**. Two columns carry the whole pipeline:

COLUMN YOU NEED	REAL NAME IN V3	WHAT IT HOLDS	USED BY
title	<code>Title</code>	the course name	part of the search text + shown in recommendations
skills	<code>Skills</code>	comma-separated skills the course teaches	the search text (M4/M5) <i>and</i> the relevance check (M6)

Step 1 — rename the v3 columns to the names the rest of the notebook expects (the real names are capitalised, so this rename matters):

```
# map the v3 dataset's column names → the names this project expects
courses_raw = courses_raw.rename(columns={
    "Title": "title",
    "Skills": "skills",
    "course_description": "description", # optional (see note)
    "Organization": "provider", # optional
    "course_url": "url", # optional
    "Difficulty": "level", # optional
})
```

Important — we do NOT use the description. v3 has a `course_description` column, but only ~65% of rows have one and some are off-topic. So the search text is built from `title` + `skills` (both clean and fully populated), *not* from the description. (Details in M2.) `Skills` is the column that matters — it feeds both the search and the scoring.

First, get the file into Colab. The CSV does not appear by itself — upload it using the  **Files panel** on the left of Colab (drag the CSV in), then read it with

```
pd.read_csv("coursera_course_dataset_v3.csv")
```

The file disappears when the Colab session resets, so re-upload if needed (or load it from your shared Google Drive / GitHub repo).

STARTER CODE

```
# 1a. the course catalogue (the v3 file from the Kaggle dataset)
courses_raw = pd.read_csv("coursera_course_dataset_v3.csv")
print(courses_raw.shape)
print(courses_raw.columns.tolist())  # LOOK at the real column names
courses_raw.head()

# 1b. role -> required skills. Hand-made is completely fine to start.
role_skills = {
    "Data Scientist": ["python", "statistics", "machine learning", "sql", "data visuali
    "Data Analyst":   ["excel", "sql", "statistics", "data visualization", "python"],
    "ML Engineer":    ["python", "machine learning", "deep learning", "sql", "cloud"],
    # ... add ~8-10 roles total
}
```

Week-1 reality check: the v3 column names are capitalised (`Title` , `Skills`) — that's exactly why the Step-1 rename matters. Always `print(courses_raw.columns.tolist())` and confirm the names before relying on them — never assume. Report the actual columns at Meeting 2.

 **For your report:** record which dataset you chose and *why*, its source URL, how many rows it has, and its real column names. List the roles in your `role_skills` table. Feeds the *Data* chapter.

✓ **Done when:** `courses_raw.head()` shows real rows, and `role_skills["Data Scientist"]` returns a list of skills.

MODULE 2 Clean & prepare

 Sprint 2 • Mtg 3

Purpose. Turn the messy raw catalogue into one tidy table with the exact columns the rest of the notebook expects. The most important output is the `text` column — the single block of words each recommender will match against.

 **Before you start M2:** this module assumes your columns are already called `title` and `skills` — you renamed them in **M1**. If you skipped that rename, the very first line `courses["title"]` will crash with `KeyError: 'title'`. Fix it by going back to M1's rename step — don't patch it here.

WHAT M2 DOES — FIVE SMALL STEPS

1. **Rename** the columns to `title` / `skills` (*done in M1 — see the callout above*).

2. **Drop junk rows** — courses missing a title or skills, and duplicate titles.
3. **Renumber** `course_id` to the clean row position (0, 1, 2, ...).
4. **Build** `text` — glue `title + skills` into one cleaned, lowercase string. *This is the field the recommenders search.* (We use skills, not the description — see M1's note on why.)
5. **Make** `skills` **a list** — split the skills cell into clean lowercase tags.

INPUT → OUTPUT

IN	OUT: COURSES — COLUMNS GUARANTEED
<code>courses_raw</code>	<code>course_id</code> (int) · <code>title</code> (str) · <code>skills</code> (list[str]) · <code>text</code> (str, cleaned = title + skills)

STARTER CODE

```
courses = courses_raw.copy()

# keep only rows that actually have a title + skills
courses = courses.dropna(subset=["title", "skills"]).drop_duplicates("title").reset_index()
courses["course_id"] = courses.index

def clean(s):
    s = re.sub(r"[^a-z0-9 ]", " ", str(s).lower())
    return re.sub(r"\s+", " ", s).strip()

# the words each recommender will match on = title + skills (NOT description)
courses["text"] = (courses["title"] + " " + courses["skills"]).apply(clean)

# NOW turn skills into a clean lowercase list (after text is built from it)
courses["skills"] = courses["skills"].fillna("").apply(
    lambda s: [x.strip().lower() for x in str(s).split(",") if x.strip()])
```

What `clean()` **does, in plain words** (the two lines look scary but are simple):

- `re.sub(r"[^a-z0-9 ]", " ", str(s).lower())` — first lower-case everything, then **replace anything that is NOT a letter, number, or space with a space** (so it strips punctuation like `/ , & ( ) : .`).
- `re.sub(r"\s+", " ", s).strip()` — **squeeze any run of spaces down to a single space**, and trim the ends.

Result: `"SQL for Data-Analysis (2024)!"` → `"sql for data analysis 2024"`. A tidy lowercase string the recommenders can match on.

⚠ **The skills split assumes COMMAS.** The line `str(s).split(",")` only works if your dataset separates tags with commas. If it uses `|` or `;` (some do), change the `","` to match — otherwise the whole cell becomes *one giant wrong "skill"* and `precision@k` (M6) breaks silently. Check with `print(courses["skills"].iloc[0])` — you should see a clean list like `['sql']`, not `['sql | python | regression']`.

WORKED EXAMPLE — ONE REAL COURSE ROW THROUGH M2

A real v3 course — *"Learn SQL Basics for Data Science"* — going from raw to clean. This is exactly what your code produces:

	BEFORE — IN <code>COURSES_RAW</code> (AFTER THE M1 RENAME)	AFTER — IN <code>COURSES</code>
<code>course_id</code>	(its original row index)	a clean row number 0...N-1
<code>title</code>	Learn SQL Basics for Data Science	Learn SQL Basics for Data Science (<i>unchanged</i>)
<code>skills</code>	"Databases, SQL, Data Management, Data Analysis, Big Data, ..." (a string)	["databases", "sql", "data management", "data analysis", "big data", ...] (a list)
<code>text</code>	<i>does not exist yet</i>	"learn sql basics for data science databases sql data management data analysis big data data model exploratory data analysis ..."

How the `text` value was made: glue `title` + a space + the `skills` string, then run `clean()` (lowercase, strip the commas and punctuation). Note `title` and `skills` aren't themselves changed — they're the raw material for `text`. When a learner's gap contains `"sql"`, this course scores high **because "sql" is in its `text`** — that is the entire matching mechanism. (*We build `text` from skills, not description, because v3's description is sparse/noisy — see M1.*)

Note on the v3 skills: they're comma-separated *phrases* (e.g. `"Statistics for Data Science"`, `"Probability & Statistics"`), which the split above turns into clean lowercase tags. That phrasing is exactly why M6's relevance check uses **whole-word matching**, not exact equality.

📄 **For your report:** document every cleaning decision — rows dropped, how you built the `text` and `skills` columns, and how many courses survived (before → after counts). Reviewers reward honest data hygiene. Feeds the *Data* chapter.

✓ **Done when:** `courses.columns` contains `course_id`, `title`, `skills`, `text`, and `courses["text"].iloc[0]` is clean lowercase words.

MODULE 3 User input + skill gap

📅 Sprint 2 · Mtg 3

Purpose. Take the two things the user provides (skills they have + the role they want), look up what that role needs, and compute what is *missing*. The gap — not the skills they have — is what we search on.

INPUT → OUTPUT

IN	OUT
<code>current_skills</code> (list) · <code>target_role</code> (str) · <code>role_skills</code>	<code>skill_gap</code> (list) · <code>query_text</code> (str)

STARTER CODE

```
# the only two things the user provides
current_skills = ["python"]
target_role    = "Data Scientist"

# Rule 2: the role must exist — fail loudly if not
assert target_role in role_skills, f"Pick a role from {list(role_skills)}"

have      = {s.lower() for s in current_skills}
# the skills still missing, kept in the role table's order (foundational first)
skill_gap = [s for s in role_skills[target_role] if s.lower() not in have]
query_text = " ".join(skill_gap)          # the search query, as one string

print("Skill gap:", skill_gap)
```

Reading that code, line by line (a few bits are new if you only know basic Python):

- `assert target_role in role_skills, "..."` — **stop immediately with a clear message** if the role isn't in the table. This is "fail loudly" — far better than the code limping on and giving empty results (Rule 2).
- `have = {s.lower() for s in current_skills}` — this `{ }` builds a **set** (a bag of unique items) of the user's skills, lower-cased. A set is used because checking "is X in here?" is fast and case is normalised.
- `skill_gap = [s for s in role_skills[target_role] if s.lower() not in have]` — read it as: **"go through each skill the role needs; keep it only if the user doesn't already have it."** What's left is the gap.
- `" ".join(skill_gap)` — glue the gap skills into one space-separated string, because the recommenders search on a single block of text.

RIYA'S WORKED EXAMPLE

NEEDED (DATA SCIENTIST)	HAS	→ SKILL_GAP
python, statistics, machine learning, sql, data visualization, deep learning	python	statistics, machine learning, sql, data visualization, deep learning

```
query_text = "statistics machine learning sql data visualization deep learning"
```

 **For your report:** explain the skill-gap idea in your own words and walk through Riya's example (needed – has = gap). State which roles your system supports. Feeds the *Methodology* chapter.

✓ **Done when:** for Riya, `skill_gap` has 5 items and `query_text` is the 5 skills joined by spaces.

MODULE 4 TF-IDF recommender (the baseline)

 Sprint 1→2 · Mtg 2→3

Purpose. The simple, honest baseline. Turn every course's `text` into a word-frequency vector, turn the `query_text` into the same kind of vector, and rank courses by how similar they are. No machine learning training — just counting and weighting words.

NEW CONCEPT — TF-IDF & COSINE SIMILARITY (READ THIS FIRST)

TF-IDF turns a piece of text into a list of numbers (a "vector") by counting its words — but cleverly: a word counts for *more* if it is rare across all courses (distinctive), and *less* if it shows up everywhere (like "course" or "learn").

Tiny example — three one-line courses:

COURSE	ITS WORDS
C1	python · machine · learning
C2	python · data · analysis
C3	yoga · meditation

`python` appears in 2 of 3 courses → common → *low* weight. `yoga` appears in only 1 → rare → *high* weight. So TF-IDF gives each course a set of numbers that emphasise its **distinctive** words.

Cosine similarity then scores how alike two of these number-vectors are, from `0` (nothing in common) to `1` (identical). We turn the user's skill gap into the same kind of vector and rank every course by its cosine similarity to it.

For the query `"machine learning"`: **C1** (python machine learning) shares both words → high score (≈ 0.8) → recommended. **C3** (yoga meditation) shares nothing → score 0 → not recommended. That ranking is exactly what this module outputs.

INPUT → OUTPUT

IN	OUT: TFIDF_RECS
<code>courses.text</code> · <code>query_text</code>	top-K DataFrame: <code>course_id</code> , <code>title</code> , <code>score</code> , <code>rank</code>

STARTER CODE

```
vec      = TfidfVectorizer(stop_words="english")
matrix   = vec.fit_transform(courses["text"])      # every course → a vector
q_vec    = vec.transform([query_text])            # the gap → same vector space

scores   = cosine_similarity(q_vec, matrix).flatten() # one score per course

K        = 10
top      = scores.argsort()[::-1][:K]              # indexes of the K best
tfidf_recs = courses.loc[top, ["course_id", "title"]].copy()
tfidf_recs["score"] = scores[top]
tfidf_recs["rank"]  = range(1, len(tfidf_recs)+1)
tfidf_recs
```

The one cryptic line, decoded: `scores.argsort()[::-1][:K]`

- `scores.argsort()` → gives the course *positions* sorted from **lowest score to highest** (worst → best).
- `[::-1]` → **reverses** it, so now it's best → worst.
- `[:K]` → takes the **first K**, i.e. the top-K courses.

So the whole line means "**give me the row positions of the K highest-scoring courses.**" (*The same trick appears in M5 — you'll recognise it.*)

Imports: `TfidfVectorizer` and `cosine_similarity` are imported once in the **M0 setup cell** — you don't re-import them here. If you get `NameError: TfidfVectorizer is not defined`, you simply haven't run the M0 cell yet.

Common error: calling `vec.fit_transform()` again on the query. Use `vec.transform()` for the query — the vectorizer was already *fitted* on the courses, and the query must live in that same vocabulary.

 **For your report:** explain TF-IDF + cosine similarity in plain English, and paste a sample top-10 table for one user (e.g. Riya). Feeds *Methodology (baseline)*.

✓ **Done when:** `tfidf_recs` has exactly K rows, ranks 1...K, and scores descending.

MODULE 5 Embedding recommender (the smart version)

 Sprint 2 · Mtg 3

Purpose. The same idea as Module 4 — rank courses by similarity to the gap — but using a pretrained model that understands *meaning*, not just exact words. So "ML" and "machine learning" land close together, where TF-IDF sees two unrelated words.

NEW CONCEPT — EMBEDDINGS (READ THIS FIRST)

Like TF-IDF, **embeddings** turn text into numbers — but instead of *counting words*, a pretrained AI model *reads* the text and produces numbers that capture its **meaning**. Texts that mean similar things get similar numbers, even when they use completely different words.

Why that matters — the same query, `"machine learning"`, against two courses:

COURSE	TF-IDF (EXACT WORDS)	EMBEDDINGS (MEANING)
"machine learning fundamentals"	✓ high — shares the words	✓ high
"deep neural networks & AI"	✗ ≈ 0 — no shared words	✓ high — related meaning

TF-IDF *misses* the second course because the words don't overlap; embeddings *catch* it because the model knows the topics are related. **Whether that actually gives better recommendations is exactly what you measure in Module 6.**

Notice: the ranking logic (cosine similarity $\rightarrow$ argsort $\rightarrow$ top-K) is *identical* to Module 4. **The only thing that changes is how text becomes numbers.** Make this point in your report — it is the cleanest way to explain the comparison.

INPUT $\rightarrow$ OUTPUT

IN	OUT: <code>EMBED_RECS</code>
<code>courses.text</code> <code>query_text</code>	same shape as <code>tfidf_recs</code> : <code>course_id</code> , <code>title</code> , <code>score</code> , <code>rank</code>

STARTER CODE

```
from sentence_transformers import SentenceTransformer
model = SentenceTransformer("all-MiniLM-L6-v2") # small, free, runs in Colab

course_emb = model.encode(courses["text"].tolist(), show_progress_bar=True)
q_emb      = model.encode([query_text])

scores = cosine_similarity(q_emb, course_emb).flatten() # SAME ranking as M4

K = 10
top = scores.argsort()[::-1][:K]
embed_recs = courses.loc[top, ["course_id", "title"]].copy()
embed_recs["score"] = scores[top]
embed_recs["rank"] = range(1, len(embed_recs)+1)
embed_recs
```

First run is slow (it downloads the model, ~90 MB, and encodes every course). That is normal. Run it once; Colab keeps it in memory after that. Keep `K` the same value as Module 4 so the comparison is fair.

 **For your report:** explain how embeddings differ from TF-IDF (meaning vs exact words), and show the *same* user's top-10 side by side with Module 4 so the contrast is visible. Feeds *Methodology (embeddings)*.

✓ **Done when:** `embed_recs` has the same columns and K rows as `tfidf_recs`, and the two lists are *different* (if identical, something is wrong).

MODULE 6 Evaluate & compare ★

 Sprint 3 · Mtg 4

Purpose. This is the graded heart of the project. Measure how good each recommender is using **precision@k** — of the top-K courses it recommended, how many were genuinely relevant (per the **locked** `is_relevant()` **rule**) — then compare, chart, and declare a winner.

NEW CONCEPT — WHAT IS PRECISION@K? (READ THIS FIRST)

You have probably seen *precision* in theory. **precision@k** is the version used for ranked lists like recommendations. Plain definition:

precision@k = out of the top k courses the recommender suggested, the fraction that were actually relevant.

Worked example. Say the recommender returns its **top 5** courses for Riya (so `k = 5`). We check each one: does it teach a skill Riya is missing?

RANK	RECOMMENDED COURSE	TEACHES A MISSING SKILL?
1	SQL for Data Analysis	✓ yes (sql)
2	Intro to Machine Learning	✓ yes (machine learning)
3	Yoga for Beginners	✗ no
4	Deep Learning Specialization	✓ yes (deep learning)
5	Cooking 101	✗ no

3 of the 5 were relevant, so:

`precision@5 = 3 / 5 = 0.6`

Higher is better. A precision@5 of 1.0 means every one of the top 5 was useful; 0.0 means none were. You compute this same number for *both* recommenders and compare — the higher one wins.

The one question this raises: who decides the ✓/✗ in that table? **You do — with one rule,** fixed in [§03](#): a course counts as relevant if one of the gap skills appears (as a whole word) in its skills. That rule is `is_relevant()`, and it is the yardstick behind every score here.

INPUT → OUTPUT

IN	OUT
<code>tfidf_recs</code> · <code>embed_recs</code> · <code>courses</code> · <code>skill_gap</code>	<code>eval_table</code> (DataFrame) · a bar chart · <code>winner</code> (str)

STARTER CODE

```
def is_relevant(course_skills, skill_gap):
    skills_text = " ".join(course_skills).lower()
    return any(re.search(r"\b" + re.escape(g) + r"\b", skills_text) for g in skill_gap)

def precision_at_k(recs, courses, skill_gap, k=10):
    hits = 0
    for cid in recs.head(k)["course_id"]:
        cskills = courses.loc[courses["course_id"]==cid, "skills"].iloc[0]
        hits += is_relevant(cskills, skill_gap)
    return hits / k

p_tfidf = precision_at_k(tfidf_recs, courses, skill_gap, k=10)
p_embed = precision_at_k(embed_recs, courses, skill_gap, k=10)

eval_table = pd.DataFrame({"method": ["TF-IDF", "Embedding"],
                           "precision@10": [p_tfidf, p_embed]})
winner = "Embedding" if p_embed >= p_tfidf else "TF-IDF"

eval_table.plot.bar(x="method", y="precision@10", legend=False)
plt.title("Which recommender is better?"); plt.ylabel("precision@10"); plt.show()
```

This is the most important code in the project — so here it is line by line:

- `re.search(r"\b" + g + r"\b", skills_text)` — checks whether the gap skill `g` appears as a **whole word** in the course's skills text; `any(... for g in skill_gap)` asks **"does at least one gap skill show up?"** → that's what makes a course "relevant". (Whole-word, so `statistics` matches `"statistics for data science"` but `sql` won't falsely match `"mysql"`.)
- `hits += is_relevant(...)` — `is_relevant` returns `True / False`, and in Python **True counts as 1, False as 0**. So this adds 1 for every relevant course. After the loop, `hits` = how many of the top-k were relevant.
- `courses.loc[courses["course_id"]==cid, "skills"].iloc[0]` — looks dense, reads simply: **"find the row whose course_id is cid, and take its skills list."** The `.iloc[0]` just grabs that single value out of the one-row result.
- `hits / k` — relevant-count ÷ how-many-we-checked = **precision@k**. Exactly the 3/5 = 0.6 from the example above.

Imports: `pd` (pandas) and `plt` (matplotlib) come from the **M0 setup cell**.

`is_relevant()` is the *same* rule locked in §03 — define it once and reuse it; don't write a second version.

Keep `k` here equal to the `K` you chose in M4/M5. If M4 returns only 5 courses but you call `precision_at_k(..., k=10)`, it divides by 10 while checking 5 — the score comes out wrongly low and the comparison is unfair.

Go further for more marks: compute precision@k for several values of k (3, 5, 10), test several different roles (not just Data Scientist), and average. **This matters here:** on this dataset a single role can score 1.0 at k=10 for *both* methods (a tie) — so testing several roles and k values is what actually reveals the difference between TF-IDF and embeddings. *This is where your ambition belongs.*

📄 **For your report — this is your *Results* chapter:** paste the `eval_table`, the bar chart, and 2–3 sentences interpreting *which method won and why*. Critically discuss the limits of your `is_relevant()` rule. This section carries the most marks — invest here.

✅ **Done when:** `eval_table` shows two precision scores, the bar chart renders, and `winner` holds one of the two method names.

MODULE 7 Pick best + LLM roadmap (Groq)

📅 Sprint 4 • Mtg 5

Purpose. Take the winning recommender's top courses and ask a free LLM to turn the bare list into an encouraging, human, month-by-month roadmap. This is the "flourish" — one API call, on top of a study that already stands on its own.

INPUT → OUTPUT

IN	OUT
<code>winner · tfidf_recs / embed_recs · current_skills · target_role</code>	<code>roadmap_text</code> (str)

STARTER CODE

```
from groq import Groq
client = Groq(api_key=os.environ["GROQ_API_KEY"])

best_recs = embed_recs if winner=="Embedding" else tfidf_recs
course_list = "\n".join(f"- {t}" for t in best_recs["title"].head(8))

prompt = f"""You are helping a learner plan their upskilling.

Current skills: {' · '.join(current_skills)}
Target role: {target_role}
Skills they still need (the gap): {' · '.join(skill_gap)}

Recommended courses — use ONLY these, do not invent or add any:
{course_list}

Write a month-by-month learning roadmap that:
- uses ONLY the courses listed above,
- orders them sensibly: foundational skills (e.g. SQL, statistics) before
  advanced ones (e.g. machine learning, deep learning),
- gives one short sentence per course on why it sits at that point,
- is realistic for a part-time learner over ~6 months,
- ends with one line stating this is a SUGGESTED order, not a verified
  prerequisite plan.

Keep the tone encouraging and concise."""

resp = client.chat.completions.create(
    model="llama-3.3-70b-versatile",          # verify current models at console.groq.com
    messages=[{"role": "user", "content": prompt}])

roadmap_text = resp.choices[0].message.content
print(roadmap_text)
```

Auth error / `KeyError: 'GROQ_API_KEY'` ? Your key isn't loaded. It's set once in the **M0 setup cell** (`os.environ["GROQ_API_KEY"] = "gsk_..."`) — re-run M0, then this cell. Never paste the key directly here, and never commit it to GitHub.

Model names change. If you get a "model not found" error, the model name was retired — check the current free models at `console.groq.com/docs/models` and swap it in. Do not guess.

Limitation to own in your report: the roadmap's *ordering* is the LLM's best guess from general knowledge — the system does **not** model real course prerequisites (no dataset carries that). Present it as a *suggested* order, not a verified one. Note in your report that true prerequisite sequencing (a prerequisite graph + topological sort) is *future work*. Also: the "use ONLY these courses" wording in the prompt is what stops the LLM inventing courses — keep it.

 **For your report:** include one full example roadmap (Riya's), and 2–3 sentences on what the LLM adds and its limits (it can't verify the courses; it just narrates). Feeds the *Conclusion / extensions* section (Student 5).

 **Done when:** `roadmap_text` is a multi-paragraph roadmap that names courses from `best_recs`. If the call fails, the project still passes — this step is optional polish.

MODULE 8

Optional UI — only if everything else is done

 Sprint 4–5 • optional

Purpose. Wrap Modules 3→7 behind simple input boxes so a non-coder can use it. **Do not start this until the notebook is complete and the report is drafted.** It earns no extra marks on its own — it is a demo nicety.

The Handbook calls this the *optional demo*. Two ways to build it: `ipywidgets` (simplest — stays inside the notebook, shown below) or a `streamlit` app (separate file, more work). Either is fine; both are optional.

SIMPLEST VERSION — INSIDE THE NOTEBOOK, NO EXTRA TOOLS

```
import ipywidgets as widgets
skills = widgets.Text(description="Your skills:") # comma-separated
role   = widgets.Dropdown(options=list(role_skills), description="Target role:")
go     = widgets.Button(description="Get my roadmap")
display(skills, role, go)
# on click: read the boxes, then re-run Modules 3 → 7 with those values
```

Remember Rule 2: the role is a `Dropdown` (fixed choices), never a free text box — otherwise an unknown role silently returns nothing.

 **For your report:** if you built it, add a screenshot and a line on how it works. If you didn't, list it as *future work* in the Conclusion — that is a perfectly good outcome. (Student 5)

✓ **Done when:** (only if attempted) typing skills + picking a role + clicking the button prints recommendations and the roadmap.

06

Plain-English glossary

TERM	IN ONE SENTENCE
TF-IDF	A way to turn text into numbers by counting words, giving rare-but-meaningful words more weight than common ones.
Embedding	A way to turn text into numbers using a pretrained model, so that pieces of text with similar <i>meaning</i> get similar numbers.
Cosine similarity	A score from 0 to 1 measuring how close two vectors point in the same direction — i.e. how similar two pieces of text are.
Skill gap	The skills a target role needs minus the skills the person already has = what they still need to learn.
precision@k	Out of the top k courses recommended, the fraction that were genuinely relevant. Higher = better recommender.
Ground truth	The "correct answer" you measure against. Here: our <code>is_relevant()</code> rule defines it.
Baseline	The simplest reasonable method (TF-IDF), used as a yardstick to judge whether the fancier method (embeddings) is actually worth it.